

ARTIFICIAL NEURAL NETWORK

in Python LANGUAGE

Chapter 4: Back Propagation & Optimization

4.3. Optimization – Other methods

- **RMS Propagation:**

- In this method, the learning rate is adaptive and is calculated **per-parameter**, i.e. it can be variable for each parameter.
- Idea: normalize parameter updates by considering history of previous updates (**cache**)
- The **bigger the sum** of the updates is, the **smaller updates** will be in future training iterations.
- Less-frequently updated parameters will keep changing.

$$\text{cache} = \text{rho} * \text{cache} + (1-\text{rho}) * \text{gradient} ** 2$$

$$\text{parameter} += \text{current_learning_rate} * \text{gradient} / (\text{sqrt}(\text{cache}) + \text{epsilon})$$

4.3. Optimization – Other methods

- **RMS Propagation:**

```
class Optimizer_RMSProp:
```

```
    def __init__(self, learning_rate = 0.1, decay = 0., epsilon = 1e-7, rho = 0.5):
```

```
        self.learning_rate = learning_rate
```

```
        self.current_learning_rate = learning_rate
```

```
        self.decay = decay
```

```
        self.step = 0
```

```
        self.epsilon = epsilon
```

```
        self.rho= rho
```

```
# pre update
```

```
def pre_update_params(self):
```

```
    if self.decay:
```

```
        self.current_learning_rate = self.learning_rate * (1 / (1 + self.decay * self.step))
```

4.3. Optimization – Other methods

- **RMS Propagation:**

Update parameters

```
def update_params(self, layer):
```

```
    if not hasattr(layer, 'weights_cache'):
```

```
        layer.weights_cache = np.zeros_like(layer.weights)
```

```
        layer.biases_cache = np.zeros_like(layer.biases)
```

```
    layer.weights_cache = self.rho * layer.weights_cache + (1-self.rho) * layer.dweights**2
```

```
    layer.biases_cache = self.rho * layer.biases_cache + (1-self.rho) * layer.dbiases**2
```

```
    layer.weights += -self.current_learning_rate * layer.dweights /  
(np.sqrt(layer.weights_cache) + self.epsilon)  
    layer.biases += -self.current_learning_rate * layer.dbiases / (np.sqrt(layer.biases_cache)  
+ self.epsilon)
```

post update

```
def post_update_params(self):
```

```
    self.step += 1
```

4.3. Optimization – Other methods

- **Adam (Adaptive momentum):**

- Currently the most widely used optimizer, combining the **momentum** concept of SGD and the **adaptive learning rate** concept of RMSProp.
- Use a **correction mechanism** to apply to the **cache** and **momentum** by dividing by the term $(1-\beta^{iterations}) \rightarrow$ This term is initially very small (important change) and will tend to 1 (stable).

4.3. Optimization – Other methods

- Adam (Adaptive momentum):

```
class Optimizer_Adam:
```

```
    def __init__(self, learning_rate = 0.1, decay = 0., epsilon = 1e-7, beta1 = 0.9, beta2 = 0.9):
        self.learning_rate = learning_rate
        self.current_learning_rate = learning_rate
        self.decay = decay
        self.step = 0
        self.epsilon = epsilon
        self.beta1 = beta1
        self.beta2 = beta2

    # pre update
    def pre_update_params(self):
        if self.decay:
            self.current_learning_rate = self.learning_rate * (1 / (1 + self.decay * self.step))
```

4.3. Optimization – Other methods

- Adam (Adaptive momentum):

```
# Update parameters
```

```
def update_params(self, layer):
```

```
    if not hasattr(layer, 'weights_cache'):
```

```
        # if layers do not contain momentum, create them then fill with zeros
```

```
        layer.weights_momentum = np.zeros_like(layer.weights)
```

```
        layer.weights_cache = np.zeros_like(layer.weights)
```

```
        layer.biases_momentum = np.zeros_like(layer.biases)
```

```
        layer.biases_cache = np.zeros_like(layer.biases)
```

```
    # Update momentum with current gradient
```

```
    layer.weights_momentum = self.beta1 * layer.weights_momentum + (1 - self.beta1) *
```

```
    layer.dweights
```

```
    layer.biases_momentum = self.beta1 * layer.biases_momentum + (1 - self.beta1) *
```

```
    layer.dbiases
```

4.3. Optimization – Other methods

- **Adam (Adaptive momentum):**

```
# Correct the momentum. step must start with 1 here
```

```
weights_momentum_corrected = layer.weights_momentum / (1 - self.beta1**(self.step +1))  
biases_momentum_corrected = layer.biases_momentum / (1 - self.beta1**(self.step +1))
```

```
# update cache
```

```
layer.weights_cache = self.beta2 * layer.weights_cache + (1-self.beta2) *  
layer.dweights**2  
layer.biases_cache = self.beta2 * layer.biases_cache + (1-self.beta2) * layer.dbiases**2
```

```
# Obtain the corrected cache
```

```
weights_cache_corrected = layer.weights_cache / (1 - self.beta2**(self.step +1))  
biases_cache_corrected = layer.biases_cache / (1 - self.beta2**(self.step +1))
```

```
# Update weights and biases
```

```
layer.weights += -self.current_learning_rate * weights_momentum_corrected /  
(np.sqrt(weights_cache_corrected) + self.epsilon)  
layer.biases += -self.current_learning_rate * biases_momentum_corrected /  
(np.sqrt(biases_cache_corrected) + self.epsilon)
```


4.3. Optimization – Other methods

- Adam (Adaptive momentum):

```
# post update  
def post_update_params(self):  
    self.step += 1
```

4.3. Optimization – Other methods

- **Regularizations L1, L2:**

- The method is used to prevent overfitting.
- Idea: Calculate **penalty** to add to the loss value of the model when **large weights and biases** exist.

L1 weight regularization: $L_{1\omega} = \lambda \sum_m |\omega_m|$

L2 weight regularization: $L_{2\omega} = \lambda \sum_m \omega_m^2$

L1 bias regularization: $L_{1b} = \lambda \sum_n |b_n|$

L2 bias regularization: $L_{2b} = \lambda \sum_n b_n^2$

$$Loss = Data_Loss + L_{1\omega} + L_{1b} + L_{2\omega} + L_{2b}$$

4.3. Optimization – Other methods

- **Regularizations L1, L2:**

Gradient calculation for Backward:

$$L_{1\omega} = \lambda \sum_m |\omega_m| \Rightarrow \frac{\partial L_{1\omega}}{\partial \omega_m} = \lambda \frac{\partial |\omega_m|}{\partial \omega_m} = \begin{cases} \lambda & \text{if } \omega_m > 0 \\ -\lambda & \text{if } \omega_m < 0 \end{cases}$$

$$L_{2\omega} = \lambda \sum_n \omega_n^2 \Rightarrow \frac{\partial L_{2\omega}}{\partial \omega_n} = \lambda \frac{\partial \omega_n^2}{\partial \omega_n} = 2\lambda \omega_n$$

4.3. Optimization – Other methods

- Regularizations L1, L2:

```
class Dense_Regularization:
```

```
    def __init__(self, n_inputs, n_neurons, weights_regularizer_l1 = 0, weights_regularizer_l2 =  
0, biases_regularizer_l1 = 0, biases_regularizer_l2 = 0):  
        # Init weights and biases  
        self.weights = 0.01*np.random.randn(n_inputs,n_neurons)  
        self.biases = np.zeros((1,n_neurons))  
        # Set regularization strength  
        self.weights_regularizer_l1 = weights_regularizer_l1  
        self.weights_regularizer_l2 = weights_regularizer_l2  
        self.biases_regularizer_l1 = biases_regularizer_l1  
        self.biases_regularizer_l2 = biases_regularizer_l2  
  
        # forward pass  
    def forward(self,inputs):  
        #calculate outputs  
        self.output = np.dot(inputs,self.weights) + self.biases  
        self.inputs = inputs
```

4.3. Optimization – Other methods

- Regularizations L1, L2:

```
def backward(self, dvalues):  
    self.dweights = np.dot(self.inputs.T, dvalues)  
    self.dbiases = np.sum(dvalues, axis = 0, keepdims = True)  
  
    if (self.weights_regularizer_l1 > 0):  
        dL1 = np.ones_like(self.weights)  
        dL1[self.weights < 0] = -1  
        self.dweights += self.weights_regularizer_l1 * dL1  
    if (self.weights_regularizer_l2 > 0):  
        self.dweights += 2 * self.weights_regularizer_l2 * self.weights  
    if (self.biases_regularizer_l1 > 0):  
        dL1 = np.ones_like(self.biases)  
        dL1[self.biases < 0] = -1  
        self.dbiases += self.biases_regularizer_l1 * dL1  
    if (self.biases_regularizer_l2 > 0):  
        self.dbiases += 2 * self.biases_regularizer_l2 * self.biases  
  
    self.dinputs = np.dot(dvalues, self.weights.T)
```

4.3. Optimization – Other methods

- Regularizations L1, L2:

In the training loop:

```
loss_data = loss_function.calculate(activation3.output,y)
regularization_loss = loss_function.regularization_loss(Dense1)
regularization_loss += loss_function.regularization_loss(Dense2)
loss = loss_data + regularization_loss
```

4.3. Optimization – Other methods

- **Drop out:**

- Another method for regularization.
- Idea: Disable some neurons at each training iteration.
- Purpose: Avoiding **overfitting** (the NN too dependent on some neurons) and **co-adoption** (when a neuron depend too much on the output of other neurons).
- Force the NN to learn with a **random part of neurons**.
- Force the model to use **more neurons** for the same task, increasing the chance of learning the underlying function of the data.

4.3. Optimization – Other methods

- **Drop out:**

```
class Dense_Dropout:
```

```
    def __init__(self, rate):
```

```
        self.rate = 1 - rate # invert the rate
```

```
    # forward pass
```

```
    def forward(self, inputs):
```

```
        self.inputs = inputs
```

```
    # generate a scaled mask
```

```
    self.binary_mask = np.random.binomial(1, self.rate, size = inputs.shape) / self.rate
```

```
    # Compute the output value
```

```
    self.output = inputs * self.binary_mask
```

```
    # backward pass
```

```
    def backward(self, dvalues):
```

```
        self.dinputs = dvalues * self.binary_mask
```


Artificial Neural Network

END OF CHAPTER 4.2 – part 2